


MODULE 04 · 30 MIN

Build Faster with Best-of-N

Claude Code Bootcamp · Day 1 · Block 4 of 10



Promise

In 30 minutes you will:

- 1 Generate **3 independent candidate implementations** of the same API.
- 2 Score each on the 3-criterion rubric: **correctness, simplicity, fit**.
- 3 Commit the winner. Throw the rest away — and feel good about it.

Why this matters

- The first thing a model produces is rarely the best thing it can produce. Variance is real and exploitable.
- Best-of-N converts a 60-second wait into a measurably better artifact, *if you have a rubric*. Without a rubric you are guessing.
- This is how senior engineers operate: propose multiple solutions, justify the pick.

Concepts

- **Best-of-N (BoN):** produce N candidates, score, pick. N=3 is the sweet spot today.
- **Three-criterion rubric:**
 - *Correctness*: does it pass the manual test plan?
 - *Simplicity*: would a junior engineer maintain this on day one?
 - *Fit*: does it match `CLAUDE.md` conventions and the existing repo style?
- **Independent candidates:** each candidate gets its own prompt context. No "now improve it"; that is iteration, not BoN.

BEST-OF-N · SCORE THE CANDIDATES, PICK THE WINNER

	A	B	C
Correctness	3	5	2
Readability	4	5	3
Tests	2	5	1
	9	15	6

★ Winner: B

Live demo flow

1. Instructor authors the prompt for a Notes API.
2. Generates **candidate A** in chat 1, saves to `candidate-a/`.
3. New chat: **candidate B**. Same prompt. Saves to `candidate-b/`.
4. Same again for **candidate C**.
5. Side-by-side scoring on the rubric. Winner committed; losers shown to class to highlight the lift.

Mini project

Notes API.

- `POST /notes` — create note `{title, body}`; returns id + timestamps.
- `GET /notes` — list with optional `?q=` substring search.
- `GET /notes/:id` — fetch one or 404.
- `PATCH /notes/:id` — update title and/or body.
- `DELETE /notes/:id` — 204.
- Persisted in **SQLite** (`notes.db`).

▶ Step-by-step lab

- STEP 1 Pick your track: **Python** (FastAPI + Pydantic v2 + sqlite3) or **Node** (Hono + Zod + better-sqlite3).
- STEP 2 Create `module-04/` and three sibling folders: `candidate-a/`, `candidate-b/`, `candidate-c/`.
- STEP 3 Open three separate Claude Code chats. Paste the same prompt into each.
- STEP 4 Save each candidate into its folder.
- STEP 5 Score each candidate on the rubric below; record in `module-04/scoring.md`.
- STEP 6 Copy the winner's files into `module-04/winner/` and commit.

Suggested Claude Code prompts

GOAL

Build a small Notes API persisting to SQLite.

CONSTRAINTS

- Track A: Python 3.11 with FastAPI + Pydantic v2 + the sqlite3 stdlib module.
- Track B: TypeScript on Node 20 with Hono + Zod + better-sqlite3.
- One process. No migrations framework – initialise the schema at startup.
- HTTP status codes: 201 on create, 200 on read/update, 204 on delete, 404 on missing, 422 on invalid body.
- Timestamps in ISO 8601 UTC.

OUTPUT FORMAT

- A runnable project (single source file is fine) plus a 5-line README with the run command.

EXAMPLES

- POST /notes {"title":"a","body":"b"} → 201 {"id":1,"title":"a","body":"b","created_at":"2026-05-30T13:00:00Z","updated_at":"2026-05-30T13:00:00Z"}
- GET /notes?q=spec → 200 [matching notes]
- GET /notes/999 → 404 {"error":"not found"}

SCORING TEMPLATE – fill once per candidate

Candidate: [a|b|c]

Correctness (0–3): can I exercise all five endpoints with curl?

Simplicity (0–3): is the source single-glance readable?

Fit (0–3): does it follow CLAUDE.md conventions?

Total: __ / 9

✓ Deliverable checklist

- ☐ [] All three candidates exist in `module-04/candidate-{a,b,c}/` .
- ☐ [] `module-04/scoring.md` has scores and one-paragraph justification per candidate.
- ☐ [] `module-04/winner/` is a clean copy of the chosen candidate.
- ☐ [] You can run the winner: `curl` against all five endpoints succeeds.

✓ Definition of done

✓ Three candidates generated independently · ✓ Rubric applied honestly · ✓ Winner runs end-to-end · ✓ Losers archived, not deleted.

▶ Review checkpoint

Pair (60 s each):

STEP 1 Read the partner's `scoring.md`.

STEP 2 Challenge one score. Was the criterion actually applied, or did it shade towards "I like this one"?

STEP 3 Confirm the partner's winner runs.

Common mistakes

- Generating one candidate, then asking "now make it better" three times. That is iteration, not BoN.
- Skipping the rubric. Without it, you pick by vibe; the lift disappears.
- Picking a winner that fails the manual test plan because it is "more elegant". Correctness gates everything else.
- Throwing the losers away before scoring is recorded.

Instructor notes

- 6 / 6 / 15 / 3 split.
- Open `skills/best-of-n/SKILL.md` live to anchor the rubric.
- If running short, drop N to 2.
- Have the manual `curl` script ready as the reference test plan.

Transition to next module



We have a winner. But "looks right" is not "is right". Next we write tests, plant a bug, and make Claude debug its own code.

Next: Module 5 — Testing, Debugging & Self-Review.